

namesofgroupmembers

October 27, 2025

1 Test Notebook

This is just an (almost) empty test notebook. You can run the following code to assert that Spark is installed and working within the container.

```
[2]: from pyspark import SparkContext, SparkConf  
  
conf = SparkConf().setMaster("local").setAppName("Spark Test")  
sc = SparkContext(conf=conf)
```



```
[3]: !pip install shapely
```



```
Requirement already satisfied: shapely in /opt/conda/lib/python3.11/site-packages (2.1.2)  
Requirement already satisfied: numpy>=1.21 in /opt/conda/lib/python3.11/site-packages (from shapely) (1.24.4)
```



```
[4]: !pip install shapely
```



```
Requirement already satisfied: shapely in /opt/conda/lib/python3.11/site-packages (2.1.2)  
Requirement already satisfied: numpy>=1.21 in /opt/conda/lib/python3.11/site-packages (from shapely) (1.24.4)
```



```
[5]: # Cell 0: imports and Spark session  
from pyspark.sql import SparkSession  
from pyspark.sql import functions as F  
from pyspark.sql.window import Window  
import json  
from shapely.geometry import shape, Point  
from pyspark.sql.types import StringType, DoubleType, LongType, IntegerType  
import math  
  
spark = SparkSession.builder \  
.appName("NYC Taxi Analysis - TP1") \  
.config("spark.sql.shuffle.partitions", "200") \  
.getOrCreate()
```

```
spark.sparkContext.setLogLevel("WARN")
```

```
[6]: import os

# Cell 1: file paths (adjust for your container)
taxi_csv_path = "../data/Sample NYC Data.csv" # <-- adjust path inside the
    ↪container
borough_geojson_path = "../data/nyc-boroughs.geojson.json" # <-- adjust
output_dir = "../data/output_tp1" # where results will be written

os.makedirs(output_dir, exist_ok=True)

print(os.path.exists(taxi_csv_path), os.path.exists(borough_geojson_path), os.
    ↪path.exists(output_dir))
```

True True True

```
[7]: # Cell 2: load CSV (select only needed columns). adapt column names if sample
    ↪differs.
# Example column names:
    ↪'hack_license', 'pickup_datetime', 'dropoff_datetime', 'pickup_longitude', 'pickup_latitude', 'd
df_raw = spark.read.option("header", "true").option("inferSchema", "true").
    ↪csv(taxi_csv_path)

cols_needed = [
    "hack_license",
    "pickup_datetime",
    "dropoff_datetime",
    "pickup_longitude",
    "pickup_latitude",
    "dropoff_longitude",
    "dropoff_latitude"
]

# Check present columns
present = [c for c in cols_needed if c in df_raw.columns]
missing = [c for c in cols_needed if c not in df_raw.columns]
if missing:
    print("WARNING: missing columns (update names to match your CSV):", missing)

df = df_raw.select(*present).na.
    ↪drop(subset=["hack_license", "pickup_datetime", "dropoff_datetime"])
df = df.withColumnRenamed("hack_license", "taxi_id") # use taxi_id as unique
    ↪driver/car id
df.show(5)
```

```

-----+-----+-----+
|      taxi_id|pickup_datetime|dropoff_datetime|pickup_longitude|pickup_l
atitude|dropoff_longitude|dropoff_latitude|
-----+-----+-----+-----+
|BA96DE419E711691B...| 01-01-13 15:11| 01-01-13 15:18| -73.978165|
40.757977| -73.989838| 40.751171|
|9FD8F69F0804BDB55...| 06-01-13 00:18| 06-01-13 00:22| -74.006683|
40.731781| -73.994499| 40.750661|
|9FD8F69F0804BDB55...| 05-01-13 18:49| 05-01-13 18:54| -74.004707|
40.73777| -74.009834| 40.726002|
|51EE87E3205C985EF...| 07-01-13 23:54| 07-01-13 23:58| -73.974602|
40.759945| -73.984734| 40.759388|
|51EE87E3205C985EF...| 07-01-13 23:25| 07-01-13 23:34| -73.97625|
40.748528| -74.002586| 40.747868|
-----+-----+-----+
-----+-----+
only showing top 5 rows

```

```
[8]: # Cell 3: load GeoJSON, build list of (borough_name, polygon) sorted by borough
      ↵code and polygon area descending
with open(borough_geojson_path, "r", encoding="utf-8") as f:
    gj = json.load(f)

features = gj.get("features", [])
polys = []
for feat in features:
    props = feat.get("properties", {})
    geom = feat.get("geometry")
    bcode = props.get("boroughCode", None)    # if available
    bname = props.get("borough", props.get("boro_name", props.get("name",
      ↵"UNKNOWN")))
    shapely_poly = shape(geom)
    area = shapely_poly.area
    polys.append({"borough": bname, "code": bcode if bcode is not None else
      ↵999, "area": area, "poly": shapely_poly})

# sort: boroughCode asc (manhattan=1) and within that area desc; but easier:
      ↵sort by (code, -area)
polys_sorted = sorted(polys, key=lambda x: (x["code"], -x["area"]))

# Build a smaller broadcastable structure: list of tuples (borough, wkt) or
      ↵keep shapely polygons directly.
# We'll broadcast a list of (borough, polygon) where polygon is shapely geometry
broad_list = [(p["borough"], p["poly"]) for p in polys_sorted]
```

```

bc_polys = spark.sparkContext.broadcast(broad_list)

print(bc_polys.value[:5]) # show first 5 borough polygons
print("Number of borough polygons broadcasted:", len(bc_polys.value))

[('Manhattan', <POLYGON ((-73.926 40.878, -73.926 40.877, -73.926 40.877,
-73.926 40.877, ....>), ('Manhattan', <POLYGON ((-73.921 40.801, -73.92 40.799,
-73.917 40.798, -73.915 40.797, -7....>), ('Manhattan', <POLYGON ((-74.017
40.693, -74.015 40.693, -74.015 40.693, -74.015 40.693, ....>), ('Manhattan',
<POLYGON ((-73.942 40.769, -73.943 40.768, -73.949 40.761, -73.952 40.758,
....>), ('Manhattan', <POLYGON ((-73.907 40.876, -73.908 40.873, -73.908 40.873,
-73.909 40.872, ....>)]
Number of borough polygons broadcasted: 104

```

```
[9]: from pyspark.sql import SparkSession
spark = SparkSession.builder.getOrCreate()
spark
```

```
[9]: <pyspark.sql.session.SparkSession at 0x7e0370249c10>
```

```
[10]: !ls ..../data
```

```
nyc-boroughs.geojson.json    output_tp1  'Sample NYC Data.csv'
```

```
[11]: # Cell 4: define UDF to map lon/lat to borough. Returns 'UNKNOWN' if not found
      ↵or coords invalid.

def lonlat_to_borough(lon, lat):
    try:
        if lon is None or lat is None:
            return "UNKNOWN"
        # some datasets use strings; cast
        lon_f = float(lon)
        lat_f = float(lat)
    except Exception:
        return "UNKNOWN"
    pt = Point(lon_f, lat_f)
    for borough, poly in bc_polys.value:
        # shapely contains
        if poly.contains(pt) or poly.touches(pt):
            return borough
    return "UNKNOWN"

from pyspark.sql.functions import udf
udf_lonlat_to_borough = udf(lonlat_to_borough, StringType())

print(udf_lonlat_to_borough)
```

```
<function lonlat_to_borough at 0x7e0369c62f20>
```

```
[12]: # Cell 5: enrich dataframe with boroughs and timestamps (unix)
# create pickup_ts and dropoff_ts in seconds from epoch
df2 = df \
    .withColumn("pickup_ts", F.unix_timestamp("pickup_datetime")) \
    .withColumn("dropoff_ts", F.unix_timestamp("dropoff_datetime")) \
    .withColumn("pickup_borough", udf_lonlat_to_borough(F.
    ↪col("pickup_longitude"), F.col("pickup_latitude")))) \
    .withColumn("dropoff_borough", udf_lonlat_to_borough(F.
    ↪col("dropoff_longitude"), F.col("dropoff_latitude"))))

# compute duration_seconds
df2 = df2.withColumn("duration_s", F.col("dropoff_ts") - F.col("pickup_ts"))

df2.show(5)
```

taxi_id	pickup_datetime	dropoff_datetime	pickup_longitude	pickup_latitude	dropoff_longitude	dropoff_latitude	pickup_borough	dropoff_borough	duration_s
BA96DE419E711691...	01-01-13 15:11	01-01-13 15:18	-73.978165						
40.757977	-73.989838	40.751171	NULL	NULL					
Manhattan	Manhattan	NULL							
9FD8F69F0804BDB55...	06-01-13 00:18	06-01-13 00:22	-74.006683						
40.731781	-73.994499	40.75066	NULL	NULL					
Manhattan	Manhattan	NULL							
9FD8F69F0804BDB55...	05-01-13 18:49	05-01-13 18:54	-74.004707						
40.73777	-74.009834	40.726002	NULL	NULL					
Manhattan	NULL								
51EE87E3205C985EF...	07-01-13 23:54	07-01-13 23:58	-73.974602						
40.759945	-73.984734	40.759388	NULL	NULL					
Manhattan	Manhattan	NULL							
51EE87E3205C985EF...	07-01-13 23:25	07-01-13 23:34	-73.97625						
40.748528	-74.002586	40.747868	NULL	NULL					
Manhattan	Manhattan	NULL							

only showing top 5 rows

```
[13]: # Cell 5: enrich dataframe with boroughs and timestamps (unix)
# create pickup_ts and dropoff_ts in seconds from epoch
```

```

df2 = df \
    .withColumn(
        "pickup_ts",
        F.unix_timestamp("pickup_datetime", "dd-MM-yy HH:mm")
    ) \
    .withColumn(
        "dropoff_ts",
        F.unix_timestamp("dropoff_datetime", "dd-MM-yy HH:mm")
    ) \
    .withColumn(
        "pickup_borough",
        udf_lonlat_to_borough(F.col("pickup_longitude"), F.
        col("pickup_latitude"))
    ) \
    .withColumn(
        "dropoff_borough",
        udf_lonlat_to_borough(F.col("dropoff_longitude"), F.
        col("dropoff_latitude"))
    )

# compute duration_seconds
df2 = df2.withColumn("duration_s", F.col("dropoff_ts") - F.col("pickup_ts"))

# verify
df2.
    select("pickup_datetime","dropoff_datetime","pickup_ts","dropoff_ts","duration_s","pickup_bor
    show(5)

```

pickup_datetime	dropoff_datetime	pickup_ts	dropoff_ts	duration_s	pickup_borough	dropoff_borough
01-01-13 15:11	01-01-13 15:18	1357053060	1357053480	420	Manhattan	Manhattan
06-01-13 00:18	06-01-13 00:22	1357431480	1357431720	240	Manhattan	Manhattan
05-01-13 18:49	05-01-13 18:54	1357411740	1357412040	300	Manhattan	Manhattan
07-01-13 23:54	07-01-13 23:58	1357602840	1357603080	240	Manhattan	Manhattan
07-01-13 23:25	07-01-13 23:34	1357601100	1357601640	540	Manhattan	Manhattan

only showing top 5 rows

```
[14]: # Cell 6: Filter - remove negative durations and durations greater than 4 hours
      ↵(14400s)
MAX_DURATION_S = 4 * 3600
df_clean = df2.filter((F.col("duration_s") >= 0) & (F.col("duration_s") <=
      ↵MAX_DURATION_S))
# optional: drop rows with UNKNOWN coords, but we keep UNKNOWN for counting if
      ↵needed
print("Total records after cleaning:", df_clean.count())
```

Total records after cleaning: 99999

```
[15]: # Cell 7: Utilization computation per taxi
# For each taxi, sort by pickup_ts and compute previous dropoff_ts (lag)
w = Window.partitionBy("taxi_id").orderBy("pickup_ts")
df_sessions = df_clean \
    .withColumn("prev_dropoff_ts", F.lag("dropoff_ts").over(w)) \
    .withColumn("idle_s_raw", F.col("pickup_ts") - F.col("prev_dropoff_ts"))

# Consider idle only if prev exists and idle >0 and <= 4h (we ignore gaps > 4h
      ↵as new session)
df_sessions = df_sessions.withColumn(
    "idle_s",
    F.when((F.col("prev_dropoff_ts").isNotNull() & (F.col("idle_s_raw") > 0) &
      ↵(F.col("idle_s_raw") <= MAX_DURATION_S),
            F.col("idle_s_raw"))).otherwise(F.lit(0))
)

# sum per taxi: occupied_time (sum duration_s) and total_idle (sum idle_s)
per_taxi = df_sessions.groupBy("taxi_id").agg(
    F.sum("duration_s").alias("occupied_s"),
    F.sum("idle_s").alias("idle_s"),
    F.count("*").alias("trips_count")
)

# compute utilization: occupied / (occupied + idle). If both zero (shouldn't
      ↵happen), set to null or 0.
per_taxi = per_taxi.withColumn(
    "utilization",
    F.when((F.col("occupied_s") + F.col("idle_s")) > 0,
           F.col("occupied_s") / (F.col("occupied_s") + F.col("idle_s"))).
      ↵otherwise(F.lit(None))
)

per_taxi.cache()
```

```
per_taxi.select("taxi_id", "trips_count", "occupied_s", "idle_s", "utilization")  
    .show(10, truncate=False)
```

taxi_id	trips_count	occupied_s	idle_s	utilization
001C8AAB90AEE49F36FCAA7B4136C81A 5 0.17870722433460076	2820	12960		
0025133AD810DBE80D35FCA8BF0BCA1F 2 0.45161290322580644	1440	2400	0.375	
002C093A2CB9FD40C8C54AB5D158FC47 17 0.45161290322580644	12600	15300		
00374328FBA75FBFCA7522671250F573 1 0.45161290322580644	960	0	1.0	
00447A6197DBB329FBF764139ACA6EC4 5 0.32326283987915405	6420	13440		
0046F1E91AA13DEDE4F6EE775C6293AB 19 0.5376044568245125	11580	9960		
00567B1CBFD51DDFAC73359B09238922 21 0.5333333333333333	11520	10080		
0057CCB5BA8D29E343B3D6D275AB22D3 1 0.5333333333333333	300	0	1.0	
006114F940CB87B3ABDCE9BF6DF6FCC4 25 0.40564635958395245	16380	24000		
006313464EC98A24BB4EBC1E2419E439 21 0.26367461430575034	11280	31500		

only showing top 10 rows

```
[16]: # Cell 8: Save per_taxi results  
per_taxi_output = output_dir + "/per_taxi_utilization"  
per_taxi.coalesce(1).write.mode("overwrite").option("header","true")  
    .csv(per_taxi_output)  
print("Saved per-taxi utilization to:", per_taxi_output)
```

Saved per-taxi utilization to: ./data/output_tp1/per_taxi_utilization

```
[17]: # Cell 9: Average time to next fare per destination borough (dropoff borough)  
# For each taxi ordered by pickup_ts compute lead(pickup_ts) to know next  
# pickup; delta = next_pickup_ts - dropoff_ts (time to next fare)  
df_next = df_clean.withColumn("next_pickup_ts", F.lead("pickup_ts").over(w)) \
```

```

    .withColumn("time_to_next_s_raw", F.col("next_pickup_ts") - F.
    ↪col("dropoff_ts"))

# consider only positive and <= 4h (session continuation)
df_next = df_next.withColumn(
    "time_to_next_s",
    F.when((F.col("time_to_next_s_raw") > 0) & (F.col("time_to_next_s_raw") <= MAX_DURATION_S), F.col("time_to_next_s_raw")).otherwise(F.lit(None))
)

# Group by dropoff_borough and compute average
avg_time_to_next = df_next.filter(F.col("time_to_next_s").isNotNull()).
    ↪groupBy("dropoff_borough").agg(
        F.count("*").alias("count_samples"),
        F.avg("time_to_next_s").alias("avg_time_to_next_s"),
        F.expr("percentile_approx(time_to_next_s, 0.5)")..
        ↪alias("median_time_to_next_s")
    ).orderBy("dropoff_borough")

avg_time_to_next.show(truncate=False)
avg_time_to_next.coalesce(1).write.mode("overwrite").option("header", "true").
    ↪csv(output_dir + "/avg_time_to_next_per_dropoff_borough")

```

dropoff_borough	count_samples	avg_time_to_next_s	median_time_to_next_s
Bronx	303	2223.5643564356437	1560
Brooklyn	2593	2086.903200925569	1380
Manhattan	78228	912.8708390857494	420
Queens	4157	2676.7187875872023	2040
Staten Island	8	4710.0	3960
UNKNOWN	1943	1498.7339166237778	720

```
[18]: # Cell 10: Trips that started and ended within the same borough
same_borough = df_clean.filter(
    (F.col("pickup_borough").isNotNull()) & (F.col("dropoff_borough").
    ↪isNotNull()) &
    (F.col("pickup_borough") == F.col("dropoff_borough")) & (F.
    ↪col("pickup_borough") != "UNKNOWN"))
).groupBy("pickup_borough").agg(F.count("*").alias("num_trips_same_borough")).
    ↪orderBy("num_trips_same_borough", ascending=False)

same_borough.show(truncate=False)
same_borough.coalesce(1).write.mode("overwrite").option("header", "true").
    ↪csv(output_dir + "/trips_same_borough")
```

```
+-----+-----+
|pickup_borough|num_trips_same_borough|
+-----+-----+
|Manhattan      |83561          |
|Queens         |1396           |
|Brooklyn        |1065           |
|Bronx          |51             |
|Staten Island  |1              |
+-----+-----+
```

[19]: # Cell 11: Trips that started in one borough and ended in another

```
cross_borough = df_clean.filter(
    (F.col("pickup_borough") != F.col("dropoff_borough")) &
    (F.col("pickup_borough") != "UNKNOWN") &
    (F.col("dropoff_borough") != "UNKNOWN"))
).groupBy("pickup_borough", "dropoff_borough").agg(F.count("*").
    →alias("num_trips")).orderBy(F.desc("num_trips"))

cross_borough.show(50, truncate=False)
cross_borough.coalesce(1).write.mode("overwrite").option("header", "true").
    →csv(output_dir + "/trips_cross_borough")
```

```
+-----+-----+-----+
|pickup_borough|dropoff_borough|num_trips|
+-----+-----+-----+
|Manhattan      |Queens          |3943   |
|Queens         |Manhattan       |3698   |
|Manhattan      |Brooklyn        |1923   |
|Brooklyn        |Manhattan       |774    |
|Queens         |Brooklyn        |597    |
|Manhattan      |Bronx           |244    |
|Brooklyn        |Queens          |115    |
|Queens         |Bronx           |100    |
|Bronx          |Manhattan       |25     |
|Manhattan      |Staten Island   |9      |
|Bronx          |Queens          |2      |
|Queens         |Staten Island   |2      |
|Staten Island  |Queens          |1      |
+-----+-----+-----+
```

[20]: # Cell 12: Basic summaries and write to one summary CSV

```
summary = {
    "total_trips_cleaned": df_clean.count(),
    "total_unique_taxis": per_taxi.count(),
    "avg_trips_per_taxi": df_clean.count() / max(per_taxi.count(), 1)
```

```

}

import json, os
os.makedirs(output_dir, exist_ok=True)
with open(os.path.join(output_dir, "summary.json"), "w") as f:
    json.dump(summary, f, indent=2)

print("Summary:", summary)

```

Summary: {'total_trips_cleaned': 99999, 'total_unique_taxis': 9990, 'avg_trips_per_taxi': 10.00990990990991}

[21]: # Cell 13: OPTIONAL: show some example trips (head) to include in report
df_clean.
↳select("taxi_id", "pickup_datetime", "dropoff_datetime", "pickup_borough", "dropoff_borough", "d
↳show(20, truncate=False)

taxi_id	pickup_datetime	dropoff_datetime	pickup_borough	dropoff_borough	duration_s
BA96DE419E711691B9445D6A6307C170	01-01-13 15:11	01-01-13 15:18	Manhattan		
Manhattan 420					
9FD8F69F0804BDB5549F40E9DA1BE472	06-01-13 00:18	06-01-13 00:22	Manhattan		
Manhattan 240					
9FD8F69F0804BDB5549F40E9DA1BE472	05-01-13 18:49	05-01-13 18:54	Manhattan		
Manhattan 300					
51EE87E3205C985EF8431D850C786310	07-01-13 23:54	07-01-13 23:58	Manhattan		
Manhattan 240					
51EE87E3205C985EF8431D850C786310	07-01-13 23:25	07-01-13 23:34	Manhattan		
Manhattan 1540					
598CCE5B9C1918568DEE71F43CF26CD2	07-01-13 15:27	07-01-13 15:38	Manhattan		
Manhattan 1660					
513189AD756FF14FE670D10B92FAF04C	08-01-13 11:01	08-01-13 11:08	Manhattan		
Manhattan 420					
CCD4367B417ED6634D986F573A552A62	07-01-13 12:39	07-01-13 13:10	Manhattan		
Queens 1860					
1DA2F6543A62B8ED934771661A9D2FA0	07-01-13 18:15	07-01-13 18:20	Manhattan		
Manhattan 300					
CD2F522EEE1FF5F5A8D8B679E23576B3	07-01-13 15:33	07-01-13 15:49	Manhattan		
Manhattan 1960					
06918214E951FA0003D1CC54955C2AB0	08-01-13 13:11	08-01-13 13:19	Manhattan		
Manhattan 480					
06918214E951FA0003D1CC54955C2AB0	08-01-13 09:50	08-01-13 10:02	Manhattan		
Manhattan 720					
E949C583ECF62C8F03FDCE1484954A08	10-01-13 12:07	10-01-13 12:17	Manhattan		
Manhattan 600					

```

|93C363DDF8ED9385D65FAD07CE3F5F07|07-01-13 07:35 |07-01-13 07:46 |Manhattan
|Manhattan      |660      |
|7CE849FEF67514F080AF80D990F7EF7F|10-01-13 15:42 |10-01-13 16:04 |Manhattan
|Manhattan      |1320     |
|7CE849FEF67514F080AF80D990F7EF7F|10-01-13 14:27 |10-01-13 14:45 |Manhattan
|Manhattan      |1080     |
|351BE7D984BE17DB2FA80A748E816472|07-01-13 22:09 |07-01-13 22:19 |Manhattan
|Manhattan      |600      |
|460C3F57DD9CB2265DB75B14CD70224D|07-01-13 17:18 |07-01-13 17:20 |Manhattan
|Manhattan      |120      |
|36773E80775F26CD1158EB5450A61C79|07-01-13 06:08 |07-01-13 06:13 |Manhattan
|Manhattan      |300      |
|D2363240A9295EF570FC6069BC4F4C92|07-01-13 22:25 |07-01-13 22:36 |Manhattan
|Manhattan      |660      |
+-----+-----+-----+
+-----+-----+
only showing top 20 rows

```

```
[30]: # =====
#   VISUALIZATIONS FOR NYC TAXI DATA ANALYSIS (TP1)
# =====

import matplotlib.pyplot as plt
import pandas as pd

# --- 1 Load saved results (replace paths if needed) ---
utilization_df = spark.read.csv(f"{output_dir}/per_taxi_utilization", ↴
    header=True, inferSchema=True).toPandas()
waiting_df = spark.read.csv(f"{output_dir}/avg_time_to_next_per_dropoff_borough", header=True, inferSchema=True). ↴
    toPandas()
same_borough_df = spark.read.csv(f"{output_dir}/trips_same_borough", ↴
    header=True, inferSchema=True).toPandas()
cross_borough_df = spark.read.csv(f"{output_dir}/trips_cross_borough", ↴
    header=True, inferSchema=True).toPandas()


# -----
# 1 Utilization distribution histogram
# -----
plt.figure(figsize=(8,5))
plt.hist(utilization_df["utilization"], bins=20, color='skyblue', ↴
    edgecolor='black')
plt.title("Distribution of Taxi Utilization")
plt.xlabel("Utilization Ratio (Occupied / Total Time)")
plt.ylabel("Number of Taxis")
plt.grid(True, linestyle='--', alpha=0.6)
```


```

plt.show()

# -----
# 2 Utilization vs Idle Time (scatter)
# -----
plt.figure(figsize=(8,5))
plt.scatter(utilization_df["idle_s"], utilization_df["utilization"], alpha=0.6, color='orange')
plt.title("Taxi Utilization vs Idle Time")
plt.xlabel("Idle Time Between Trips (seconds)")
plt.ylabel("Utilization Ratio (Occupied / Total Time)")
plt.grid(True, linestyle='--', alpha=0.6)
plt.show()

```





```
[23]: # -----
# 3 Average Waiting Time by Borough (bar)
# -----
print("Columns in waiting_df:", list(waiting_df.columns)) # check real column names first

# Normalize column names (lowercase)
waiting_df.columns = [c.lower() for c in waiting_df.columns]

# Detect column for average time
time_col = None
for c in waiting_df.columns:
    if "avg" in c and "time" in c:
        time_col = c
        break

if time_col is None:
    raise KeyError("Could not find average time column in waiting_df!")

# Detect borough column
borough_col = None
for c in waiting_df.columns:
    if "boro" in c:
```

```


borough_col = c
break

if borough_col is None:
    raise KeyError("Could not find borough column in waiting_df!")

# Now plot safely
waiting_df = waiting_df.sort_values(time_col, ascending=True)
plt.figure(figsize=(8,5))
plt.bar(waiting_df[borough_col], waiting_df[time_col], color='mediumseagreen')
plt.title("Average Waiting Time Before Next Fare by Borough")
plt.xlabel("Dropoff Borough")
plt.ylabel("Average Waiting Time (seconds)")
plt.xticks(rotation=45)
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.show()

```

Columns in waiting_df: ['dropoff_borough', 'count_samples',
'avg_time_to_next_s', 'median_time_to_next_s']



```
[26]: import pandas as pd
import matplotlib.pyplot as plt

# Load the same-borough trip results from the saved CSV
same_df = pd.read_csv("../data/output_tp1/trips_same_borough/
↪part-00000-5325a43b-fa14-4809-9629-923c9c64f545-c000.csv")

print("Loaded Same-Borough Trips:")
print(same_df.head())

# Normalize column names to lowercase
same_df.columns = [c.lower() for c in same_df.columns]


# Detect column names automatically
boro_col = next((c for c in same_df.columns if "boro" in c), same_df.columns[0])
count_col = next((c for c in same_df.columns if "count" in c or "num" in c), ↪
↪same_df.columns[-1])

# Sort by trip count
same_df = same_df.sort_values(count_col, ascending=False)

# ---- Visualization ----
plt.figure(figsize=(8,5))
plt.bar(same_df[boro_col], same_df[count_col], color="cornflowerblue")
plt.title("Trips Starting and Ending in the Same Borough", fontsize=13)
plt.xlabel("Borough", fontsize=11)
plt.ylabel("Number of Trips", fontsize=11)
plt.grid(axis="y", linestyle="--", alpha=0.6)
plt.tight_layout()
plt.show()
```

Loaded Same-Borough Trips:

	pickup_borough	num_trips_same_borough
0	Manhattan	83561
1	Queens	1396
2	Brooklyn	1065
3	Bronx	51
4	Staten Island	1



```
[28]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import glob

# ---- Load the cross-borough trips CSV ----
# (your PySpark notebook saved it in ../data/output_tp1/trips_cross_borough)
csv_path = glob.glob("../data/output_tp1/trips_cross_borough/part-00000*.
˓→CSV") [0]
cross_df = pd.read_csv(csv_path)

print("Loaded Cross-Borough Trips:")
print(cross_df.head())

# ---- Normalize column names ----
cross_df.columns = [c.lower() for c in cross_df.columns]

# ---- Detect relevant columns automatically ----
pickup_col = next((c for c in cross_df.columns if "pickup" in c), None)
dropoff_col = next((c for c in cross_df.columns if "dropoff" in c), None)
count_col = next((c for c in cross_df.columns if "num" in c or "count" in c),
˓→None)

# ---- Pivot table to make matrix ----
```

```


pivot_df = cross_df.pivot_table(index=pickup_col, columns=dropoff_col,
                                values=count_col, aggfunc="sum", fill_value=0)

# ---- Heatmap Visualization ----
plt.figure(figsize=(8,6))
sns.heatmap(pivot_df, annot=True, fmt=".0f", cmap="YlGnBu", cbar_kws={'label':'Number of Trips'})
plt.title("Cross-Borough Taxi Trip Matrix", fontsize=14)
plt.xlabel("Drop-off Borough", fontsize=11)
plt.ylabel("Pickup Borough", fontsize=11)
plt.tight_layout()
plt.show()

```

Loaded Cross-Borough Trips:

	pickup_borough	dropoff_borough	num_trips
0	Manhattan	Queens	3943
1	Queens	Manhattan	3698
2	Manhattan	Brooklyn	1923
3	Brooklyn	Manhattan	774
4	Queens	Brooklyn	597



[]: